

# Data Wrangling with R: Day 1

Base R and tidyverse

Presented by Emi Tanaka

Department of Econometrics and Business Statistics



MONASH University

1st December 2020 @ Statistical Society of Australia | Zoom

# Base R

- **Base R** refers to the state of R when just launched in a clean interactive development environment (IDE, e.g. RStudio).
- In this state, R attached the following packages:
  - `base`,
  - `methods`, for complex reasons, you may need to explicitly call this when using Rscript and package development!
  - `stats`,
  - `graphics`,
  - `grDevices`,
  - `utils`, and
  - `datasets`.

# Tidyverse

- What packages does `library(tidyverse)` load?
- **Tidyverse** refers to a collection of R-packages: `ggplot2`, `dplyr`, `tidyR`, `readr`, `purrr`, `tibble`, `stringr`, `forcats`, DBI, haven, httr, readxl, rvest, jsonlite, xml2, lubridate, hms, blob, magrittr, glue and more recently, `tidymodels`.
- Eight of these packages form the **core tidyverse**.



- `library(tidyverse)` is a short hand for `library(ggplot2)`, `library(dplyr)`, ..., `library(forcats)`

# Base R and Tidyverse

- Tidyverse is not a substitute for Base R 
- Knowing Base R is essential to use Tidyverse effectively 
- Data wrangling in Tidyverse just gives you a different flavour of how you can do things in Base R 
- All data wrangling can be achieved using Base R, so why load extra package(s) to deal with the data wrangling? 
- Tidyverse packages share a common design philosophy, grammar and data structure
- This trains your mental model to do data science tasks in a certain manner which may make it easier, faster, and/or fun for you to do these tasks
- It's an opinionated approach so you should make a decision for yourself of what works for you *and* others that you share your work with

# data.frame Base R

- In R, `data.frame` is a special class of `list`
- Every column in the `data.frame` is a vector of the same length
- Each entry in a vector have the same type, e.g. logical, integer, double (or numeric), character or factor.
- It has an attribute `row.names` which could be a sequence of integers indexing the row number or some unique identifier

cars			mtcars										
##	<i>speed</i>	<i>dist</i>	##		<i>mpg</i>	<i>cyl</i>	<i>disp</i>	<i>hp</i>	<i>drat</i>	<i>wt</i>	<i>qsec</i>	<i>vs</i>	
## 1	4	2	## Mazda RX4		21.0	6	160.0	110	3.90	2.620	16.46	0	
## 2	4	10	## Mazda RX4 Wag		21.0	6	160.0	110	3.90	2.875	17.02	0	
## 3	7	4	## Datsun 710		22.8	4	108.0	93	3.85	2.320	18.61	1	
## 4	7	22	## Hornet 4 Drive		21.4	6	258.0	110	3.08	3.215	19.44	1	

# Tidy data

🎯 The typical aim of data wrangling is to make data into a **tidy data**



## Definition of a tidy data

- Each variable must have its own column
- Each observation must have its own row
- Each value must have its own cell

country	date	cases	deaths
USA	20/1/24	148,65	8,7
India	20/1/24	37,429	4,8
Brazil	20/1/24	16,603	3,4
France	20/1/24	4,552	500
Russia	20/1/24	25,73	3,1

variables

country	date	cases	deaths
USA	20/1/24	148,65	8,7
India	20/1/24	37,429	4,8
Brazil	20/1/24	16,603	3,4
France	20/1/24	4,552	500
Russia	20/1/24	25,73	3,1

observations

country	date	cases	deaths
USA	20/1/24	148,65	8,7
India	20/1/24	37,429	4,8
Brazil	20/1/24	16,603	3,4
France	20/1/24	4,552	500
Russia	20/1/24	25,73	3,1

values

# Data wrangling with Base R

# Subsetting by column

Base R Part 1

By column names

```
mtcars[, c("mpg", "cyl")]
```

```
##                                     mpg cyl
## Mazda RX4                  21.0   6
## Mazda RX4 Wag              21.0   6
## Datsun 710                 22.8   4
## Hornet 4 Drive             21.4   6
## Hornet Sportabout          18.7   8
## Valiant                     18.1   6
## Duster 360                 14.3   8
## Merc 240D                  24.4   4
## Merc 230                    22.8   4
## Merc 280                    19.2   6
## Merc 280C                  17.8   6
## Merc 450SE                 16.4   8
```

By index

```
mtcars[, 1:2]
```

```
##                                     mpg cyl
## Mazda RX4                  21.0   6
## Mazda RX4 Wag              21.0   6
## Datsun 710                 22.8   4
## Hornet 4 Drive             21.4   6
## Hornet Sportabout          18.7   8
## Valiant                     18.1   6
## Duster 360                 14.3   8
## Merc 240D                  24.4   4
## Merc 230                    22.8   4
## Merc 280                    19.2   6
## Merc 280C                  17.8   6
## Merc 450SE                 16.4   8
```

# Subsetting by column

Base R Part 2

By column names

```
mtcars[c("mpg", "cyl")]
```

```
##                                     mpg cyl
## Mazda RX4           21.0   6
## Mazda RX4 Wag       21.0   6
## Datsun 710          22.8   4
## Hornet 4 Drive      21.4   6
## Hornet Sportabout   18.7   8
## Valiant              18.1   6
## Duster 360          14.3   8
## Merc 240D            24.4   4
## Merc 230              22.8   4
## Merc 280              19.2   6
## Merc 280C             17.8   6
## Merc 450SE            16.4   8
```

By index

```
mtcars[1:2]
```

```
##                                     mpg cyl
## Mazda RX4           21.0   6
## Mazda RX4 Wag       21.0   6
## Datsun 710          22.8   4
## Hornet 4 Drive      21.4   6
## Hornet Sportabout   18.7   8
## Valiant              18.1   6
## Duster 360          14.3   8
## Merc 240D            24.4   4
## Merc 230              22.8   4
## Merc 280              19.2   6
## Merc 280C             17.8   6
## Merc 450SE            16.4   8
```

# Lists Base R

- Remember, `data.frame` is just a special type of `list` and inherit methods applied for `list`.

```
x <- list( int = 1:3,  
          shop = c("apple", "pear"),  
          c(2020, 12, 1, 1.3))
```

```
x
```

```
## $int  
## [1] 1 2 3  
##  
## $shop  
## [1] "apple" "pear"  
##  
## [[3]]  
## [1] 2020.0    12.0     1.0      1.3
```

```
x[c("int", "shop")]  
## $int  
## [1] 1 2 3  
##  
## $shop  
## [1] "apple" "pear"
```

Double square bracket to access inside the list

```
x[["int"]] # or x$int  
## [1] 1 2 3
```

# list and data.frame

Base R

```
x[, "int"] # this is a special operation for data frames  
  
## Error in x[, "int"]: incorrect number of dimensions  
  
mtcars[["cyl"]]  
  
## [1] 6 6 4 6 8 6 8 4 4 6 6 8 8 8 8 8 8 4 4 4 4 8 8 8 8 4 4 4 4 8 6 8 4  
  
mtcars$cyl  
  
## [1] 6 6 4 6 8 6 8 4 4 6 6 8 8 8 8 8 8 4 4 4 4 8 8 8 8 4 4 4 4 8 6 8 4  
  
mtcars[["cyl"]]  
  
##  
## Mazda RX4  
## Mazda RX4 Wag  
## Datsun 710
```



data.frame inherits methods for list, but list does not inherit methods for data.frame

# Subsetting by column

Base R



- If you subset a *single column* using `[ , ]`, then by default the output is a vector and not a `data.frame`.

```
mtcars[, "mpg"]
```

```
## [1] 21.0 21.0 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 17.8 16.4 17.3 15.2 10  
## [16] 10.4 14.7 32.4 30.4 33.9 21.5 15.5 15.2 13.3 19.2 27.3 26.0 30.4 15.8 19  
## [31] 15.0 21.4
```

- If you want to preserve the output as a `data.frame` then:

```
mtcars[, "mpg", drop = FALSE]
```

```
##                               mpg  
## Mazda RX4                 21.0  
## Mazda RX4 Wag              21.0  
## Datsun 710                 22.8
```

# Subsetting by row Base R Part 1

By index:

```
mtcars[3:1, ]
```

```
##          mpg cyl disp  hp drat    wt  qsec vs am gear carb
## Datsun 710 22.8   4 108  93 3.85 2.320 18.61  1  1     4     1
## Mazda RX4 Wag 21.0   6 160 110 3.90 2.875 17.02  0  1     4     4
## Mazda RX4     21.0   6 160 110 3.90 2.620 16.46  0  1     4     4
```

If it has row names:

```
mtcars[c("Datsun 710", "Mazda RX4"), ]
```

```
##          mpg cyl disp  hp drat    wt  qsec vs am gear carb
## Datsun 710 22.8   4 108  93 3.85 2.32 18.61  1  1     4     1
## Mazda RX4  21.0   6 160 110 3.90 2.62 16.46  0  1     4     4
```

# Subsetting by row Base R Part 2

Using a logical vector,

```
mtcars[mtcars$mpg > 31, ]  
  
##          mpg cyl disp hp drat      wt    qsec vs am gear carb  
## Fiat 128 32.4   4 78.7 66 4.08 2.200 19.47  1  1     4     1  
## Toyota Corolla 33.9   4 71.1 65 4.22 1.835 19.90  1  1     4     1
```

Or using **non-standard evaluation**:

```
subset(mtcars, mpg > 31)  
  
##          mpg cyl disp hp drat      wt    qsec vs am gear carb  
## Fiat 128 32.4   4 78.7 66 4.08 2.200 19.47  1  1     4     1  
## Toyota Corolla 33.9   4 71.1 65 4.22 1.835 19.90  1  1     4     1
```

Non-standard evaluation? We'll come back to this later.

# Adding or modifying a column

Base R

Append the new column:

```
df1 <- cbind(mtcars, gpm = 1 / mtcars$mpg)
```

Create a new column as you would a create a new object in a list:

```
df1$gpm <- 1 / df1$mpg  
df1[["gpm"]] <- 1 / df1$mpg
```

Overwrite column if modifying an existing column:

```
df1$wt <- df1$gpm^2  
df1[["wt"]] <- df1$gpm^2
```

Modify only a part of it:

```
df1$wt[df1$cyl==6] <- 10
```

# Adding a row Base R

What do you notice with the order of the new entry?

```
df2 <- rbind(cars, data.frame(dist = 10, speed = 3))  
tail(df2, 3)
```

```
##      speed dist  
## 49      24   120  
## 50      25    85  
## 51      3     10
```

```
df2 <- rbind(cars, c(10, 3))  
tail(df2, 3)
```

```
##      speed dist  
## 49      24   120  
## 50      25    85  
## 51      10     3
```

# Sorting columns Base R

```
mtcars[, sort(names(mtcars))]
```

	<i>am</i>	<i>carb</i>	<i>cyl</i>	<i>disp</i>	<i>drat</i>	<i>gear</i>	<i>hp</i>	<i>mpg</i>	<i>qsec</i>	<i>vs</i>	<i>wt</i>
## Mazda RX4	1	4	6	160.0	3.90	4	110	21.0	16.46	0	2.620
## Mazda RX4 Wag	1	4	6	160.0	3.90	4	110	21.0	17.02	0	2.875
## Datsun 710	1	1	4	108.0	3.85	4	93	22.8	18.61	1	2.320
## Hornet 4 Drive	0	1	6	258.0	3.08	3	110	21.4	19.44	1	3.215
## Hornet Sportabout	0	2	8	360.0	3.15	3	175	18.7	17.02	0	3.440
## Valiant	0	1	6	225.0	2.76	3	105	18.1	20.22	1	3.460
## Duster 360	0	4	8	360.0	3.21	3	245	14.3	15.84	0	3.570
## Merc 240D	0	2	4	146.7	3.69	4	62	24.4	20.00	1	3.190
## Merc 230	0	2	4	140.8	3.92	4	95	22.8	22.90	1	3.150
## Merc 280	0	4	6	167.6	3.92	4	123	19.2	18.30	1	3.440
## Merc 280C	0	4	6	167.6	3.92	4	123	17.8	18.90	1	3.440
## Merc 450SE	0	3	8	275.8	3.07	3	180	16.4	17.40	0	4.070
## Merc 450SL	0	3	8	275.8	3.07	3	180	17.3	17.60	0	3.730

# Sorting rows Base R

```
order(mtcars$mpg)
```

```
## [1] 15 16 24 7 17 31 14 23 22 29 12 13 11 6 5 10 25 30 1 2 4 32 21 3  
## [26] 8 27 26 19 28 18 20
```

```
mtcars[order(mtcars$mpg), ]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
## Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
## Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
## Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
## Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
## Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
## Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
## AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
## Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2

# Calculating statistical summaries by group

Base R

- 🎯 Calculate the *average* weight (wt) of a car for each gear type in (gear) mtcars

```
tapply(mtcars$wt, mtcars$gear, mean)
```

```
##          3          4          5  
## 3.892600 2.616667 2.632600
```

- 🎯 Calculate the *median* weight (wt) of a car for each gear (gear) and engine (vs) type in mtcars

```
tapply(mtcars$wt, list(mtcars$gear, mtcars$vs), median)
```

```
##          0          1  
## 3 3.8100 3.215  
## 4 2.7475 2.550  
## 5 2.9700 1.513
```

# Session Information

```
devtools::session_info()
```

```
## - Session info -----
##   setting  value
##   version  R version 4.0.1 (2020-06-06)
##   os        macOS Catalina 10.15.7
##   system   x86_64, darwin17.0
##   ui        X11
##   language (EN)
##   collate  en_AU.UTF-8
##   ctype    en_AU.UTF-8
##   tz       Australia/Melbourne
##   date     2020-11-26
##
## - Packages -----
##   package * version date      lib source
##   anicon     0.1.0   2020-06-21 [1] Github (emitanaka/anicon@0b756df)
##   assertthat  0.2.1   2019-03-21 [2] CRAN (R 4.0.0)
```

These slides are licensed under

